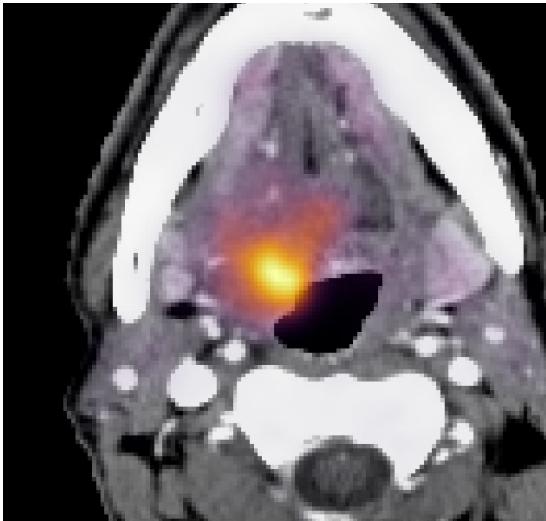
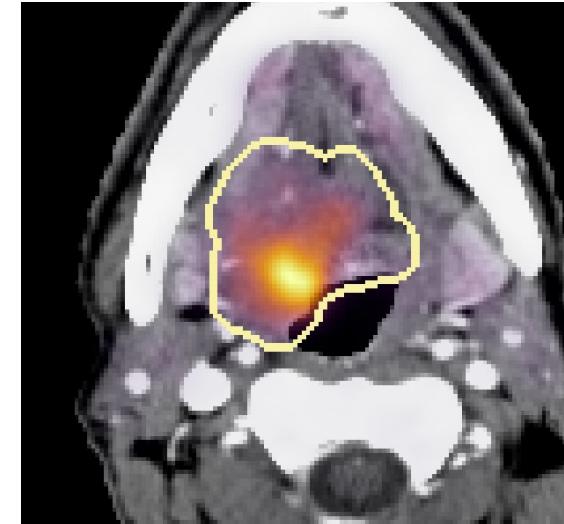


# Automatisk inntegning av kreftsvulster i medisinske bilder av pasienter med hode- og nakkekreft



Maskinlæring



**Yngve Mardal Moe**



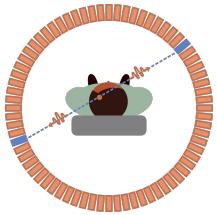
Norges miljø- og  
biovitenskapelige  
universitet



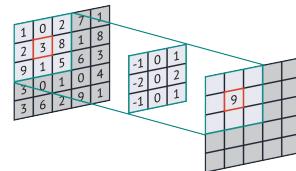
UNIVERSITETET  
I OSLO



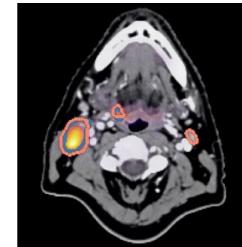
Oslo  
universitetssykehus



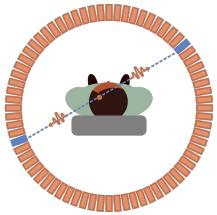
## Medisinske bilder



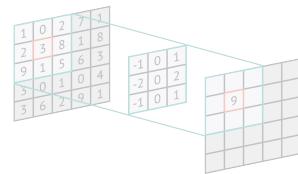
Digitale bilder, dyp læring  
og konvolusjonelle nett



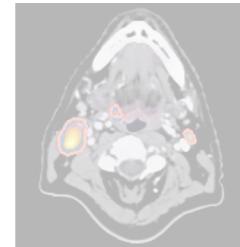
Resultater og  
observasjoner



## Medisinske bilder

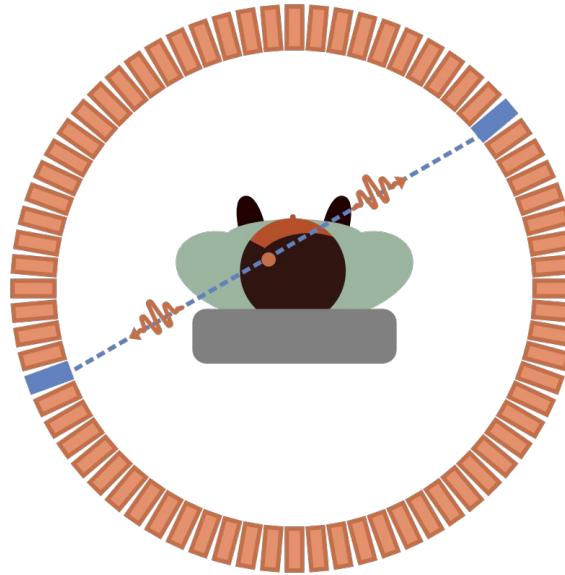
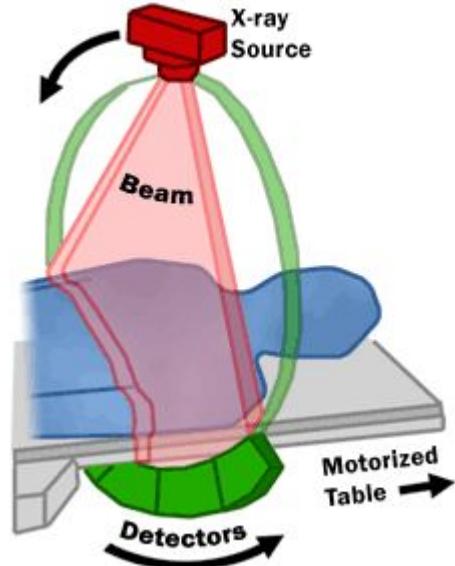


Digitale bilder, dyp læring  
og konvolusjonelle nett

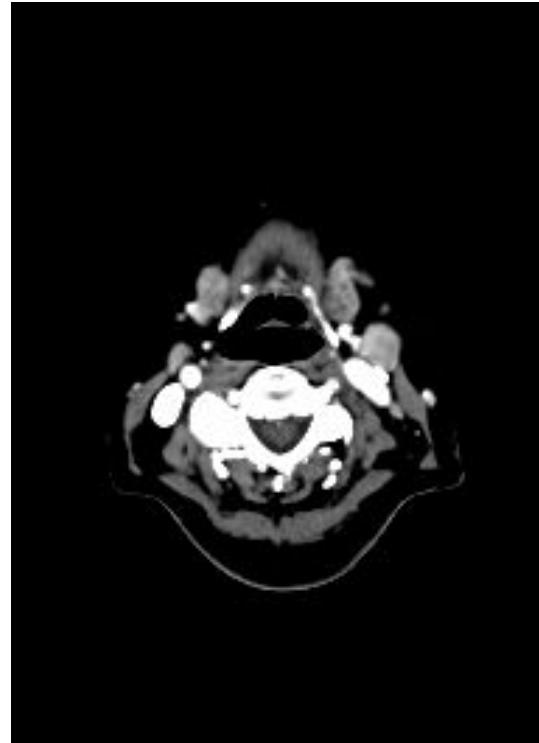
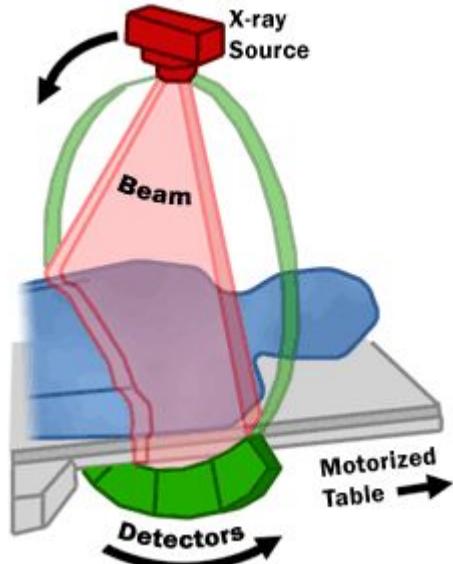


Resultater og  
observasjoner

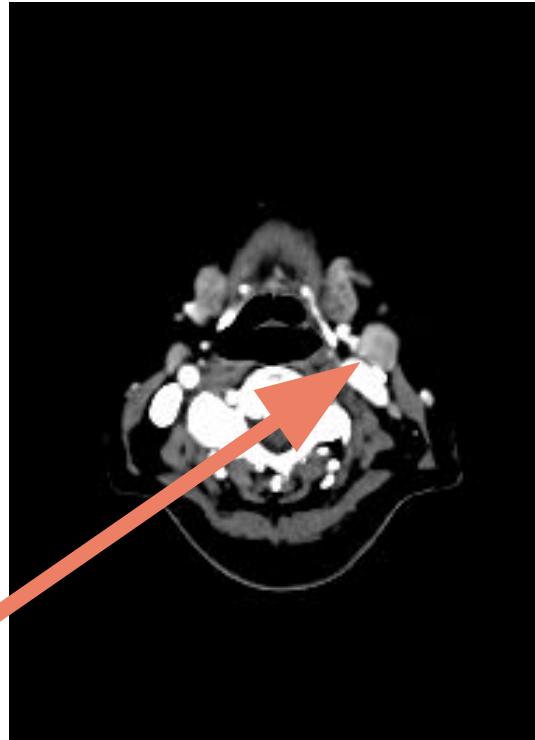
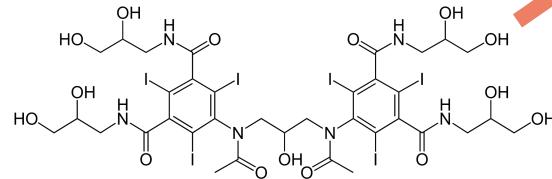
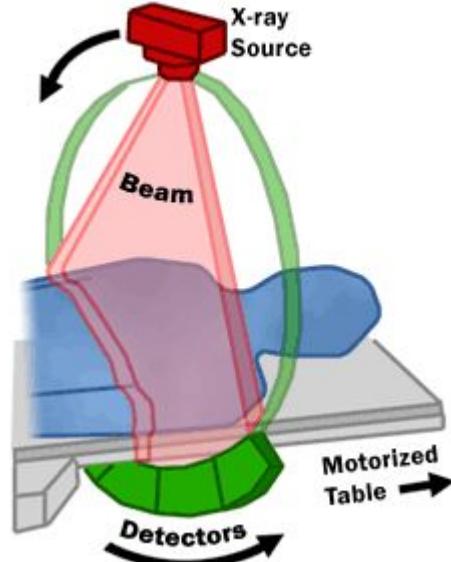
# For å finne kreftsvulster brukes en kombinasjon av PET- og CT-bilder



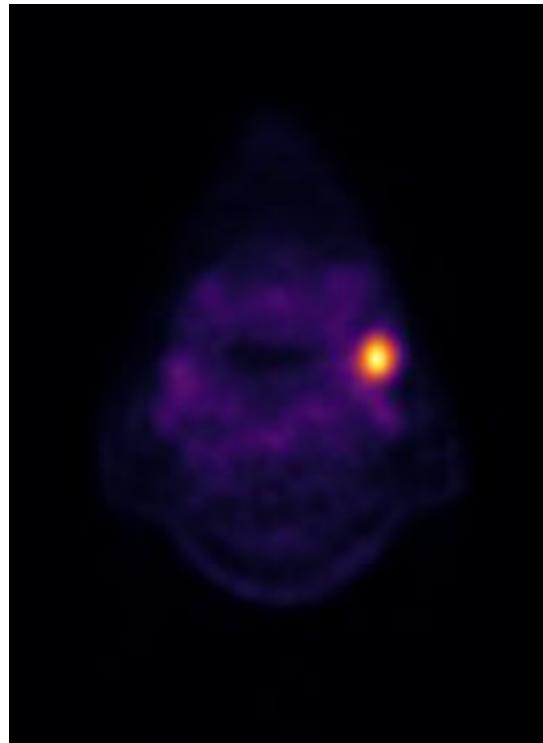
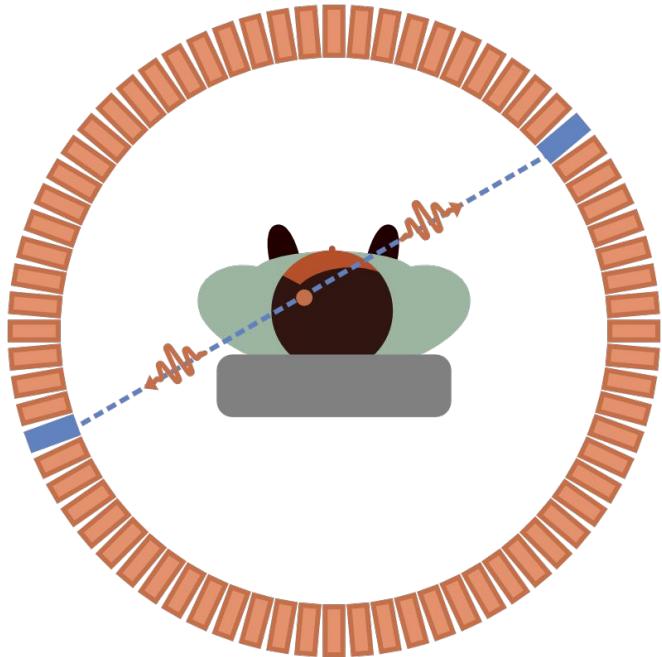
# CT bilder måler massetettheten i kroppen

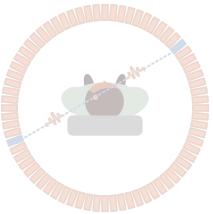


# For å gjøre svulsten tydeligere, injiseres pasientene med kontrastvæske

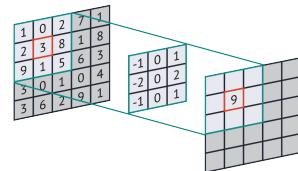


# PET-metoden som brukes gir et mål på energiforbruk i kroppen

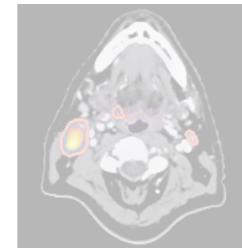




## Medisinske bilder

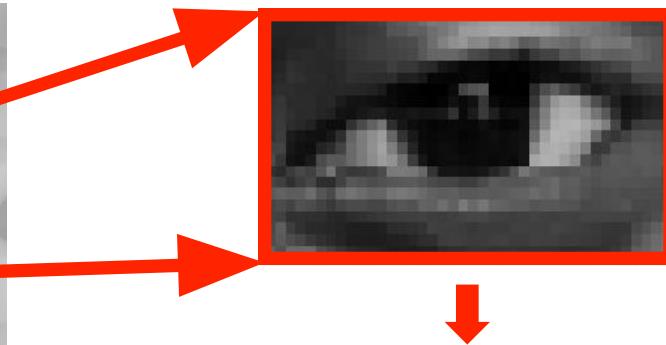
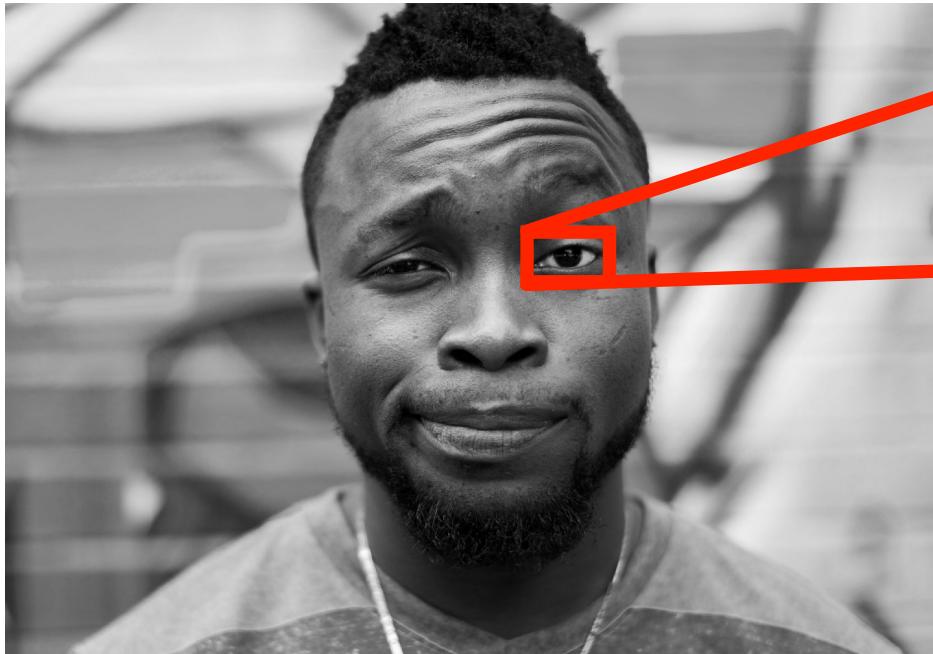


**Digitale bilder, dyp læring  
og konvolusjonelle nett**



**Resultater og  
observasjoner**

# Datamaskiner representerer digitale bilder som en mengde tall



16	11	10	16	24	40	51	61
12	12	14	19	26	58	60	55
14	13	16	24	40	57	69	56
14	17	22	29	51	87	80	62
18	22	37	56	68	109	103	77
24	35	55	64	81	104	113	92
49	64	78	87	103	121	120	101

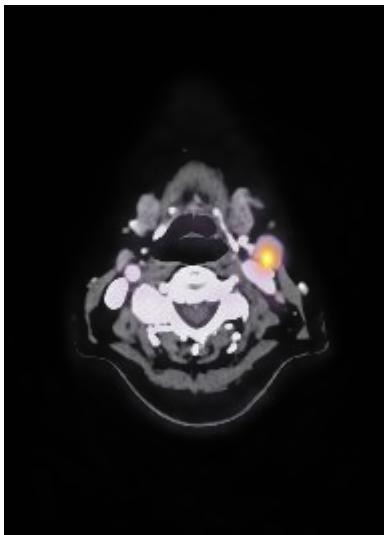
# Og fargeinformasjon i forskjellige *fargekanaler*



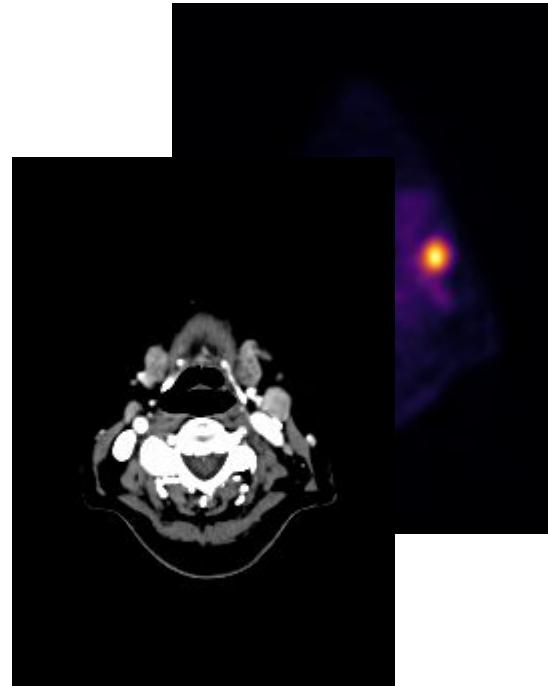
=



# Slike bildekanaler kan representere andre ting enn farger



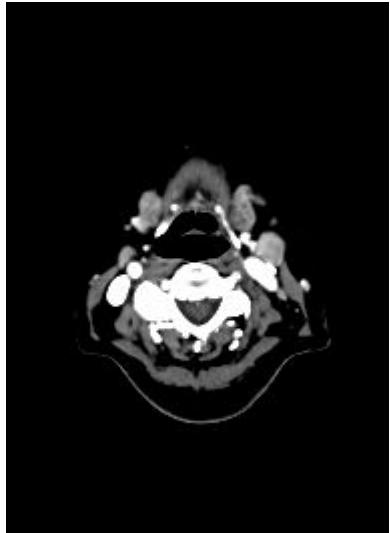
=



# Vi ønsker metoder som finner nyttig informasjon i bilder



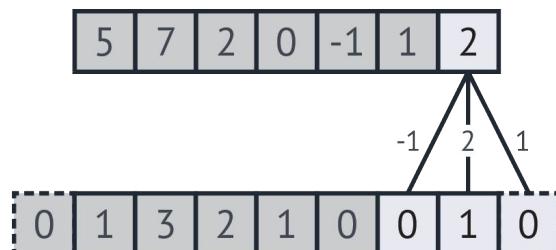
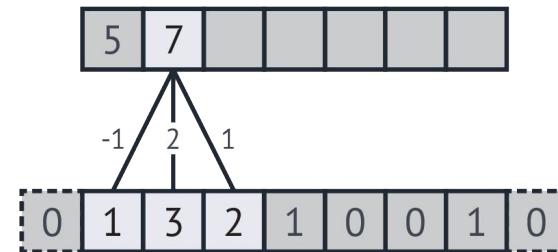
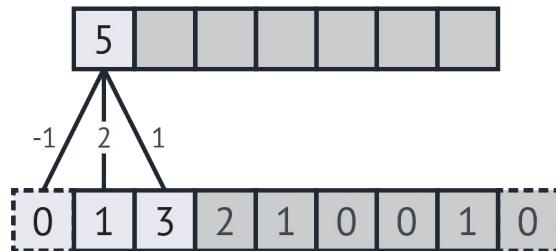
# Kanter er et eksempel på nyttig informasjon i bilder



Kantdeteksjon



# Hvordan kan datamaskiner finne informasjon i bilder?



1	0	2	7	1
2	3	8	1	8
9	1	5	6	3
3	0	1	0	4
3	6	2	9	1

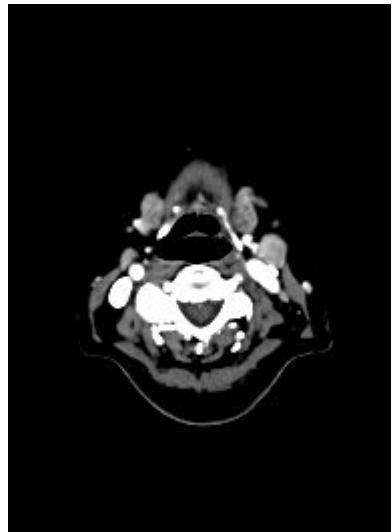
-1	0	1
-2	0	2
-1	0	1

		9	

# Konvolusjonslag består av en konvolusjon etterfulgt av terskling



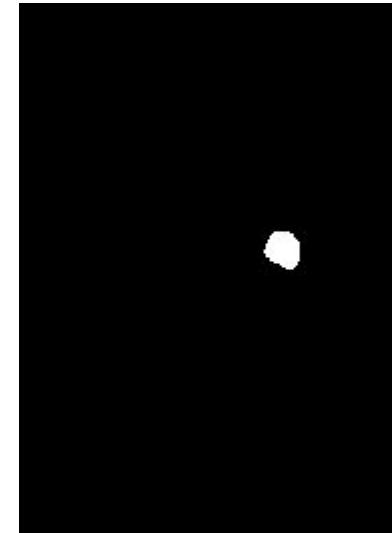
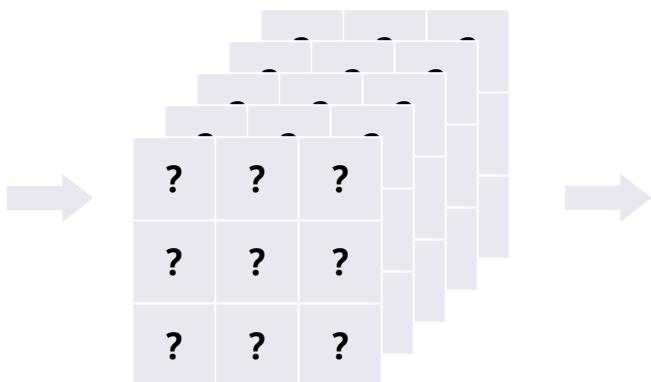
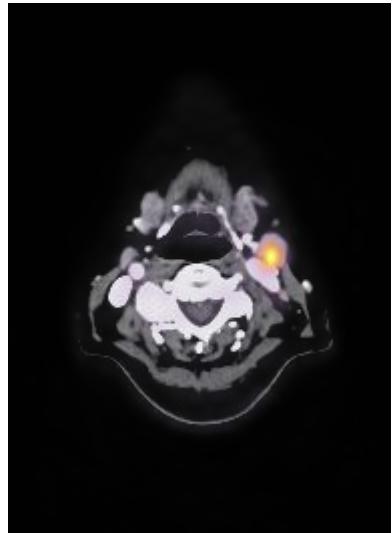
# Kanter er et eksempel på nyttig informasjon i bilder



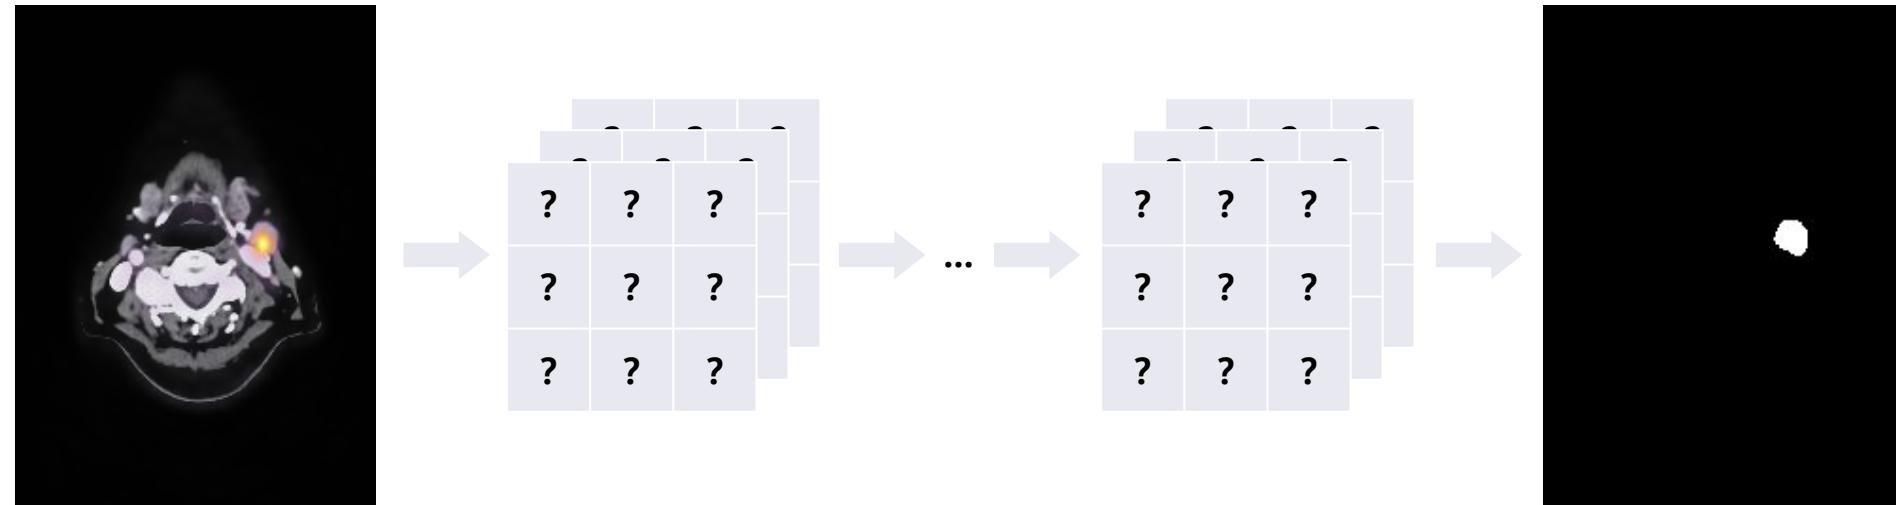
0	1	0
1	-4	1
0	1	0



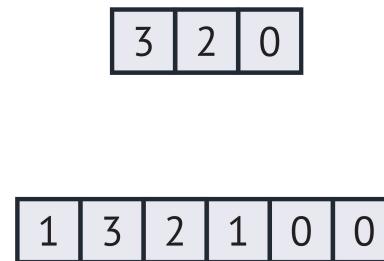
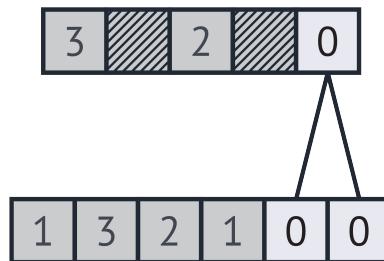
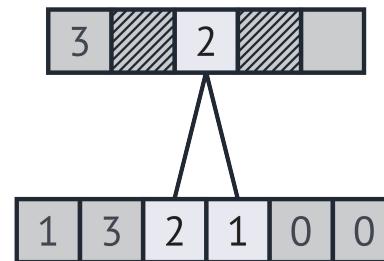
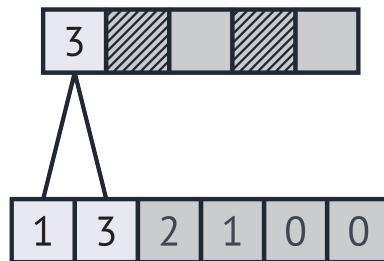
# Men vi trenger mer enn bare kanter for å finne kreft



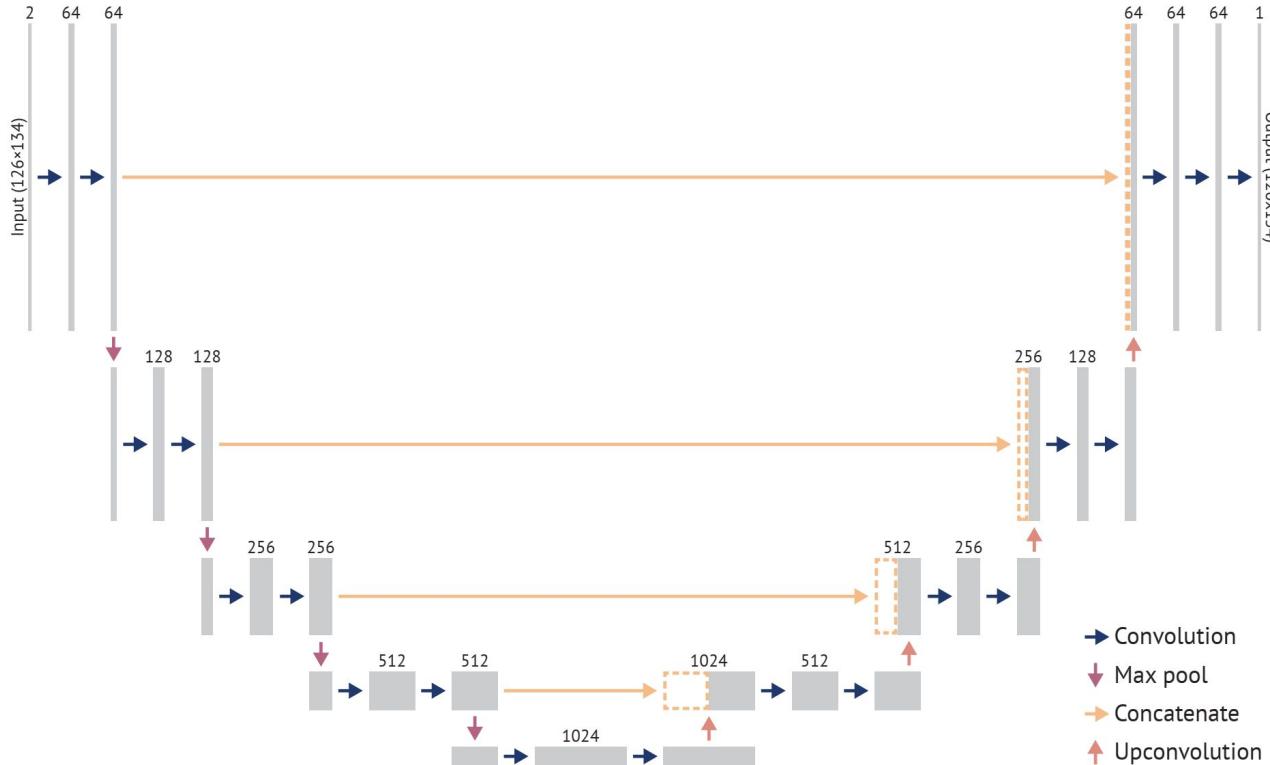
Ved å sette slike konvolusjonslag etter hverandre får vi et konvolusjonelt nett



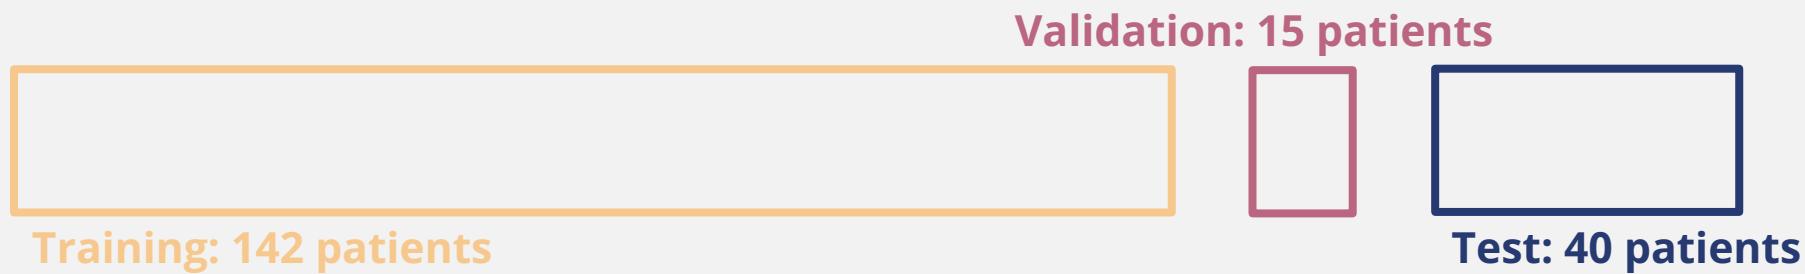
# Vi pleier også nedskalere bildene mellom konvolusjonslagene

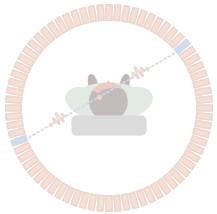


# For bildesegmentering brukes ofte U-Net modellen

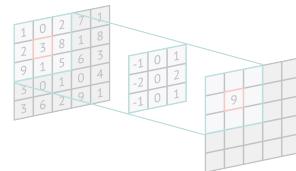


# Vi delte datasettet inn i tre deler for å evaluere modellen

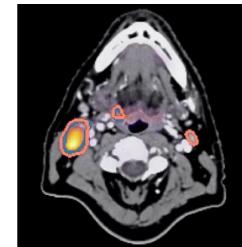




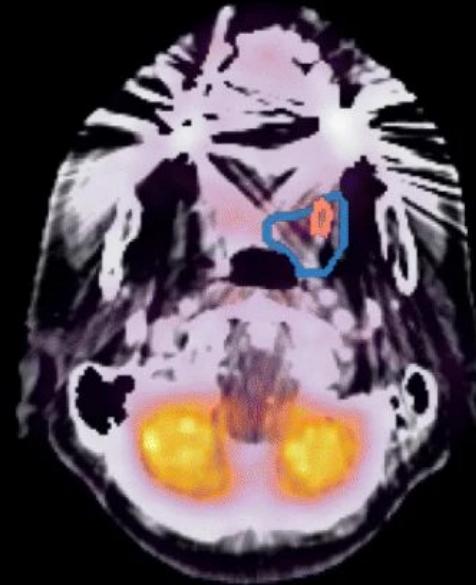
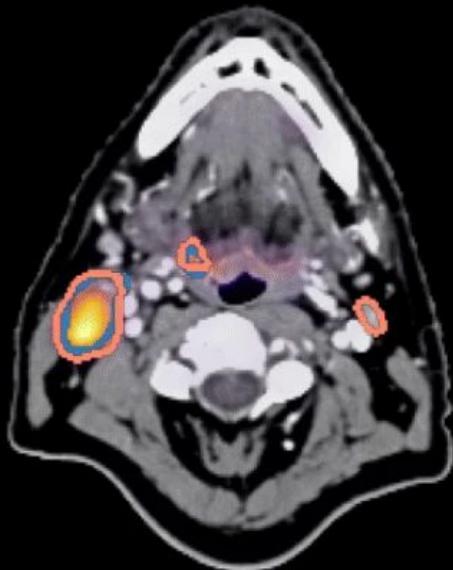
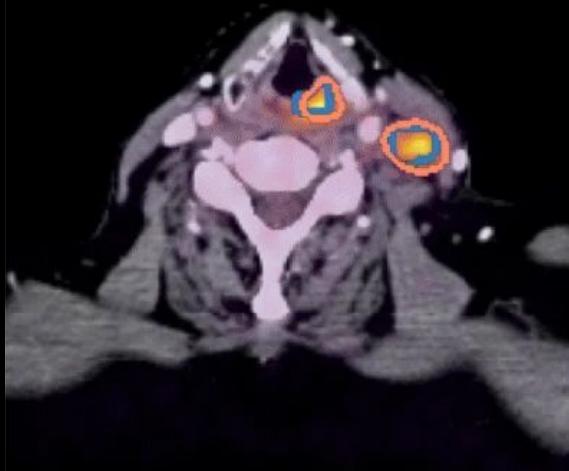
## Medisinske bilder



Digitale bilder, dyp læring  
og konvolusjonelle nett



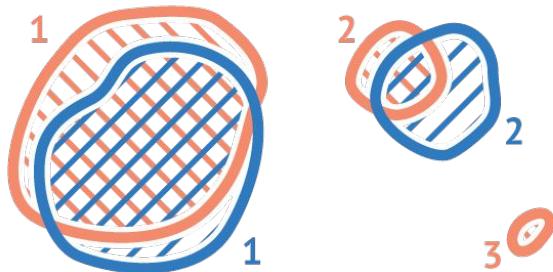
**Resultater og  
observasjoner**



CNN  
Ground truth

## Dice score

---

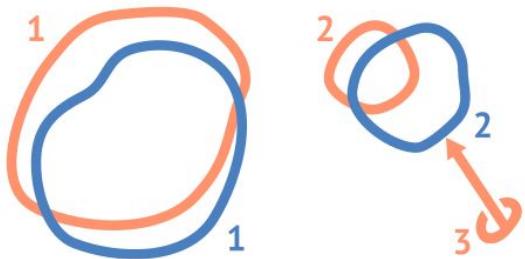


$$\text{Dice} = \frac{\frac{\square}{\square}}{\frac{\square + \square}{2}}$$

71% overlapp

## **Vanlig målt avstand**

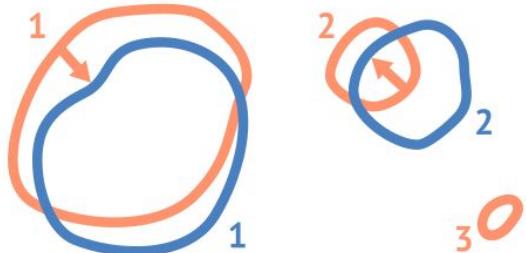
---



**4.7 mm avstand**

## **Strukturvis målt avstand**

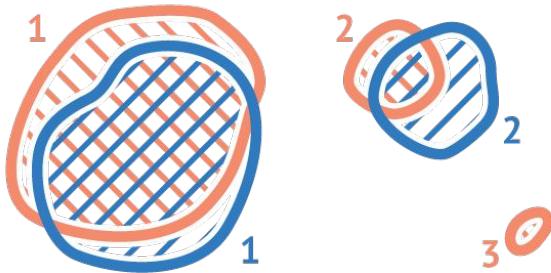
---



**1.0 mm avstand**

## Nøyaktighet

---



**86 % sensitiv**

**33 % nøyaktig**

# De falske positive inntegningene var ofte veldig små

Sanne inntegninger:

$15 \text{ cm}^3$

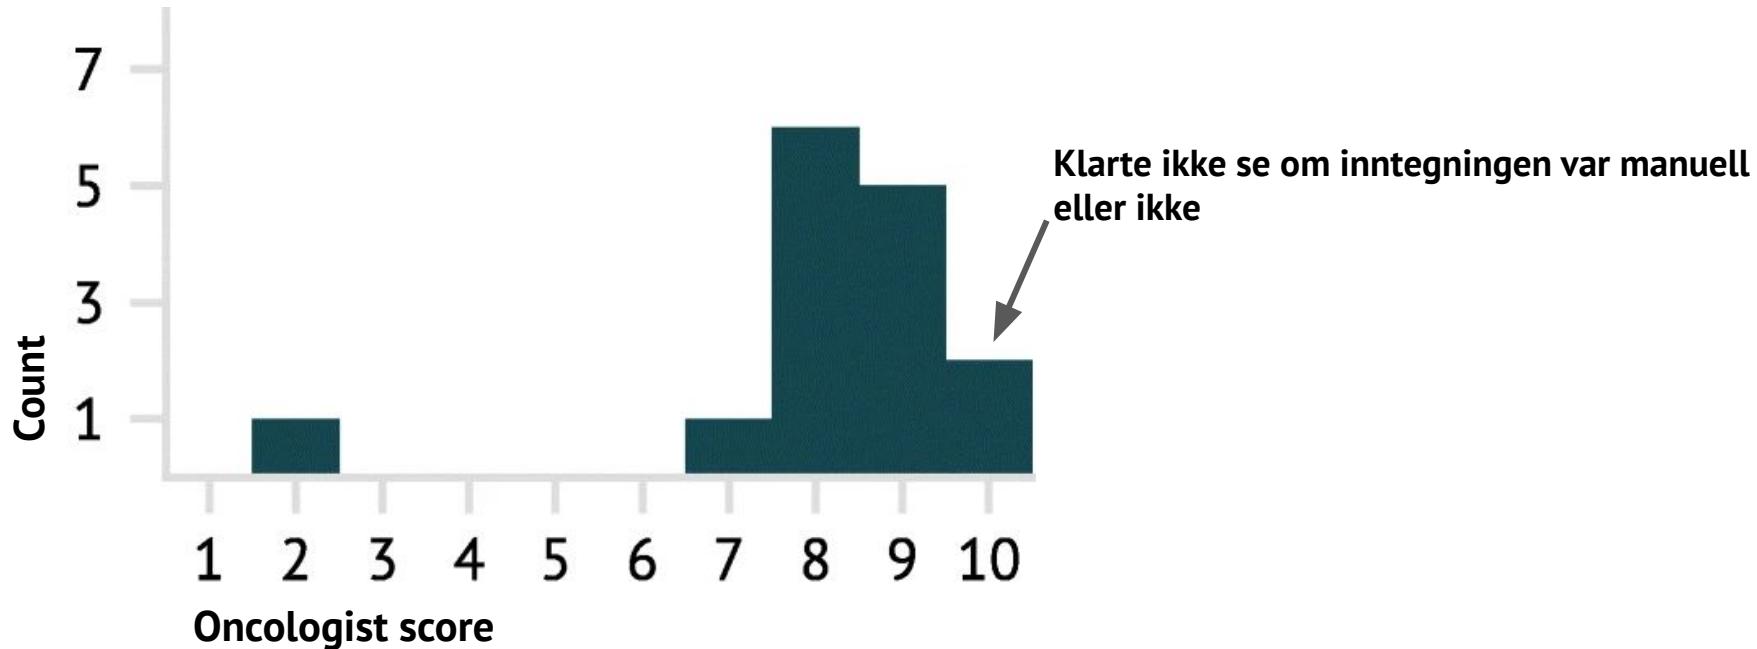
Falskt positive inntegninger:

$0.54 \text{ cm}^3$

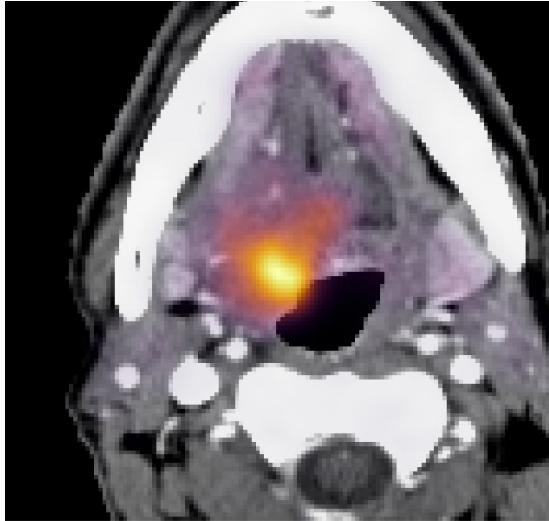
**Vi hadde også hjelp av en erfaren onkolog for å evaluere  
inntegningene**



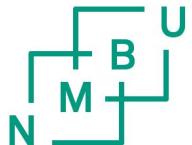
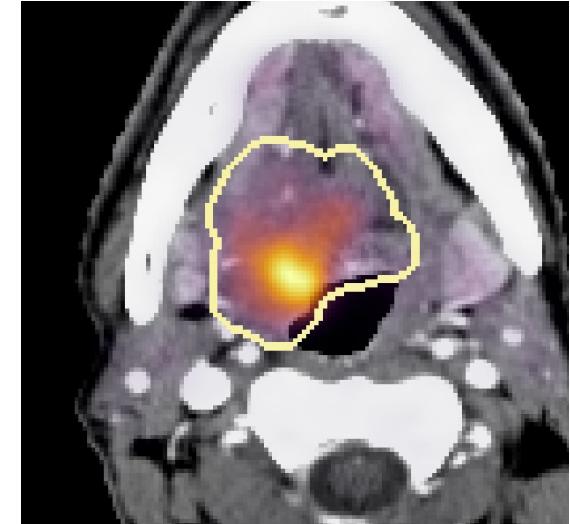
# Onkologen var generelt veldig fornøyd med auto-inntegningene



# Automatiske inntegningsmetoder kan brukes for å minke arbeidsbyrden under stråleterapiplanlegging



Maskinlæring



Norges miljø- og  
biovitenskapelige  
universitet



UNIVERSITETET  
I OSLO

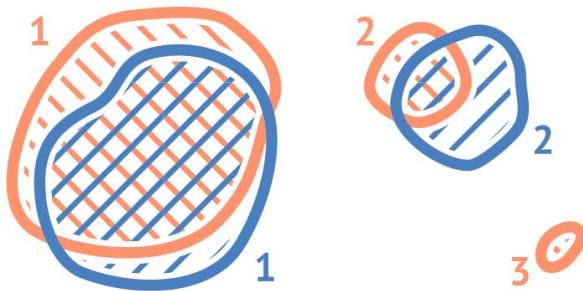


Oslo  
universitetssykehus

**Takk for meg**

# Additionally, we introduced new surface-wise performance metrics

## Coverage Fraction (CFrac)



$$CFrac_1 = \frac{\text{O}_1}{\text{S}_1} = 0.7$$

$$CFrac_2 = \frac{\text{O}_2}{\text{S}_2} = 0.6$$

$$CFrac_3 = \frac{\text{O}_3}{\text{S}_3} = 0$$

$$CFrac_1 = \frac{\text{O}_1}{\text{D}_1} = 0.8$$

$$CFrac_2 = \frac{\text{O}_2}{\text{D}_2} = 0.3$$

— CNN  
— Oncologist

## Structure-wise sensitivity and PPV

$$PPV_{CNN} = 0.7$$

$$Sens_{GT} = 0.5$$



**In deep learning, we hope to find functions that utilise  
the structure of our data to perform a given task**

$$f(x) \approx y$$

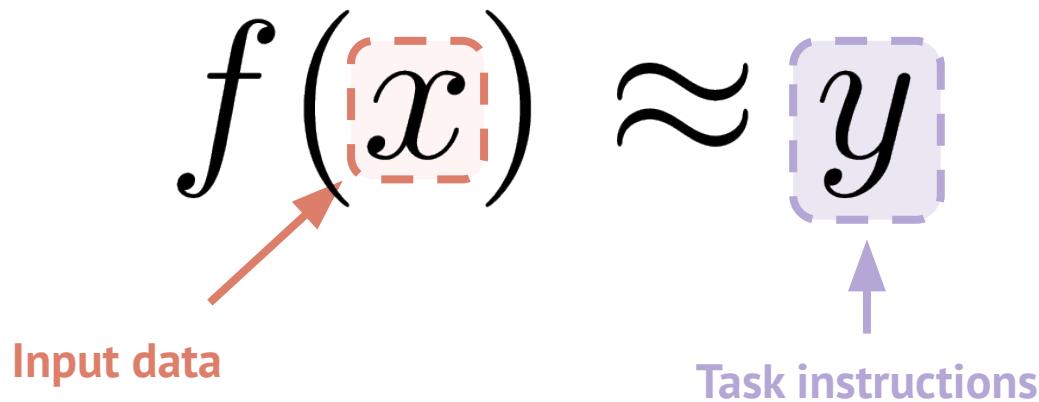
**In deep learning, we hope to find functions that utilise  
the structure of our data to perform a given task**

$$f(\tilde{x}) \approx y$$

Input data

The diagram illustrates a function  $f$  mapping input data  $\tilde{x}$  to output  $y$ . The input  $\tilde{x}$  is shown in red, with a red arrow pointing from the text 'Input data' to it. The function  $f$  is represented by a black curve. The output  $y$  is also shown in red.

In deep learning, we hope to find functions that utilise the structure of our data to perform a given task



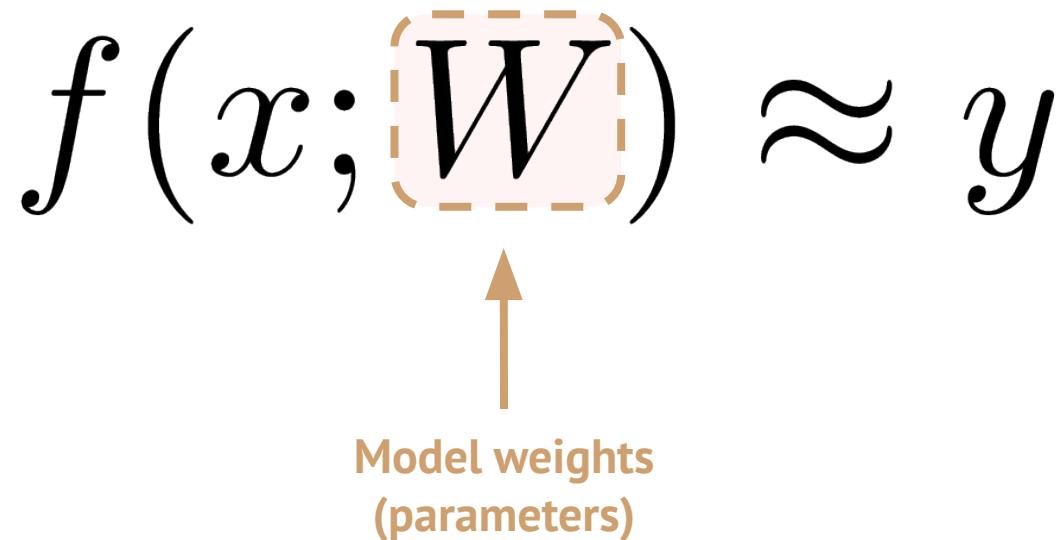
In deep learning, we hope to find functions that utilise the structure of our data to perform a given task



**To discover this function, we parametrise it by a set of weights  
and try to find the optimal weights**

$$f(x; W) \approx y$$

To discover this function, we parametrise it by a set of weights and try to find the optimal weights

$$f(x; \tilde{W}) \approx y$$


Model weights  
(parameters)

# **Neural networks are described as a cascade of function compositions**

$$f(x; W) = \check{f}_n(\check{f}_{n-1}(\dots \check{f}_2(\check{f}_1(x; W_1), W_2)\dots, W_{n-1}), W_n)$$

# Neural networks are described as a cascade of function compositions

$$f(x; W) = \check{f}_n(\check{f}_{n-1}(\dots \check{f}_2(\check{f}_1(x; W_1), W_2), \dots, W_{n-1}), W_n)$$

**Neural network**

**Layer functions**

The diagram illustrates the structure of a neural network as a composition of functions. A red dashed box labeled "Neural network" contains the mathematical expression  $f(x; W) = \check{f}_n(\check{f}_{n-1}(\dots \check{f}_2(\check{f}_1(x; W_1), W_2), \dots, W_{n-1}), W_n)$ . Inside this box, a sequence of purple dashed boxes labeled "Layer functions" encloses the nested functions  $\check{f}_1, \check{f}_2, \dots, \check{f}_n$ . Red arrows point upwards from the "Neural network" box to the first layer function  $\check{f}_1$ , and a purple arrow points upwards from the "Layer functions" label to the second layer function  $\check{f}_2$ .

**Each layer is generally an affine map (linear plus constant)  
followed by an element-wise nonlinear activation function**

$$f(x; W) = \check{f}_n(\check{f}_{n-1}(\dots \check{f}_2(\check{f}_1(x; W_1), W_2), \dots, W_{n-1}), W_n)$$

$$\check{f}_k(x; A, b) = \phi(Ax + b)$$

**Each layer is generally an affine map (linear plus constant)  
followed by an element-wise nonlinear activation function**

$$f(x; W) = \check{f}_n(\check{f}_{n-1}(\dots \check{f}_2(\check{f}_1(x; W_1), W_2), \dots, W_{n-1}), W_n)$$

$$\check{f}_k(x; \boxed{A, b}) = \phi(Ax + b)$$



Layer weights ( $W_k$  above)

**Each layer is generally an affine map (linear plus constant) followed by an element-wise nonlinear activation function**

$$f(x; W) = \check{f}_n(\check{f}_{n-1}(\dots \check{f}_2(\check{f}_1(x; W_1), W_2), \dots, W_{n-1}), W_n)$$

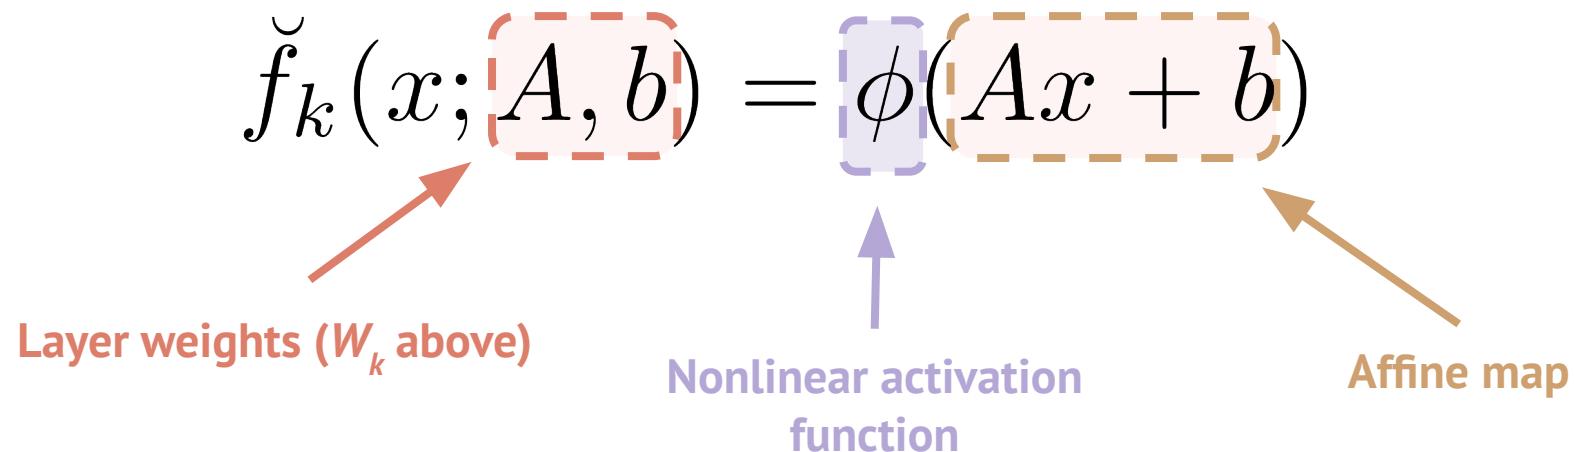
$$\check{f}_k(x; A, b) = \phi(Ax + b)$$

Layer weights ( $W_k$  above)

Affine map

**Each layer is generally an affine map (linear plus constant) followed by an element-wise nonlinear activation function**

$$f(x; W) = \check{f}_n(\check{f}_{n-1}(\dots \check{f}_2(\check{f}_1(x; W_1), W_2), \dots, W_{n-1}), W_n)$$



**The most commonly used activation function is the ReLU function**

$$\check{f}_k(x; A, b) = \phi(Ax + b)$$

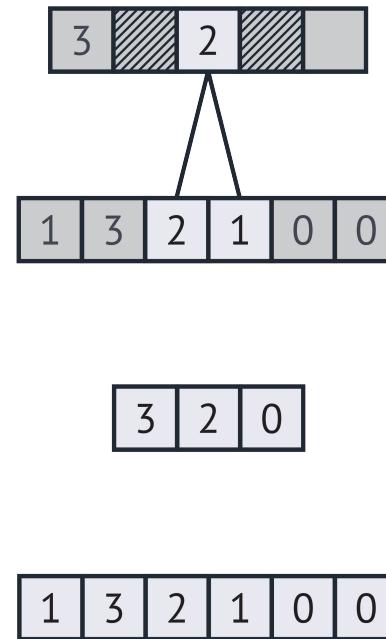
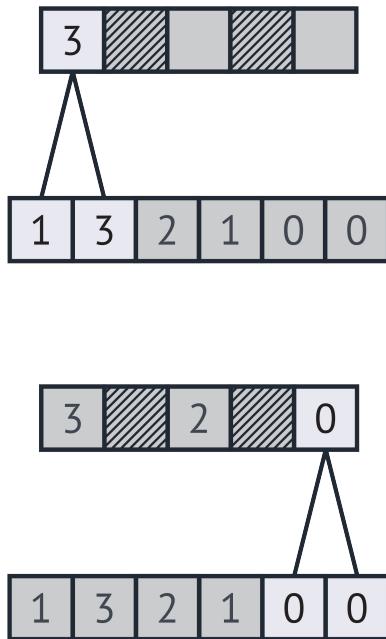
$$\phi(x) = \max(x, 0)$$



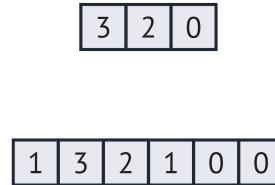
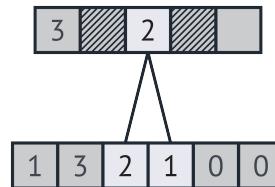
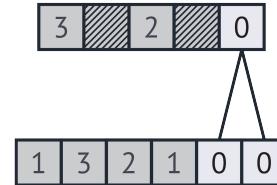
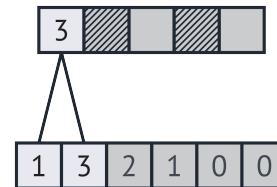
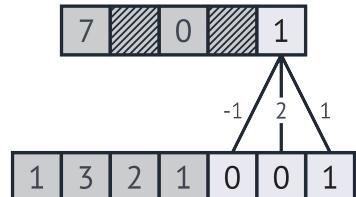
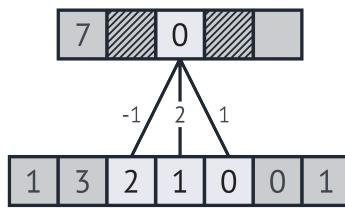
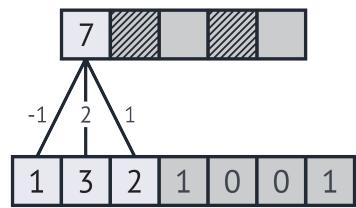
**If we just stack convolutions, then the receptive field grows slowly**



To combat this, we generally use downsampling, or *pooling* layers



# Downsampling is generally either performed via a max-pooling or a strided convolution



# This upscaling is normally done via a transposed strided convolution



↑ Downsample



↑ Conv



**Strided convolution**



↓ Upsample



↓ Conv



**Transposed strided convolution**  
“Upconvolution”  
“Deconvolution”  
[...]

**We formulate a loss function that measures the severity of our prediction errors**

$$L(W; \mathcal{T}) = \sum_{(x^{(i)}, y^{(i)}) \in \mathcal{T}} l(f(x^{(i)}; W), y^{(i)})$$

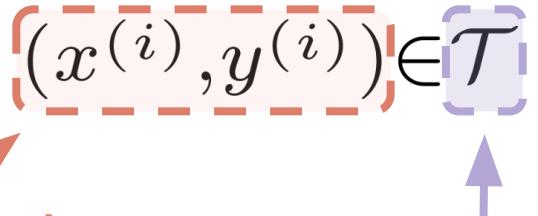
We formulate a loss function that measures the severity of our prediction errors

$$L(W; \mathcal{T}) = \sum_{(x^{(i)}, y^{(i)}) \in \mathcal{T}} l(f(x^{(i)}; W), y^{(i)})$$

Training example

We formulate a loss function that measures the severity of our prediction errors

$$L(W; \mathcal{T}) = \sum l(f(x^{(i)}; W), y^{(i)})$$



Training example

Training set

We formulate a loss function that measures the severity of our prediction errors

$$L(W; \mathcal{T}) = \sum_{\substack{((x^{(i)}, y^{(i)})) \in \mathcal{T}}} l(f(x^{(i)}; W), y^{(i)})$$

Diagram illustrating the components of the loss function formula:

- Training example:**  $(x^{(i)}, y^{(i)})$  is highlighted with a red dashed box and has a red arrow pointing to it.
- Training set:**  $\mathcal{T}$  is highlighted with a purple dashed box and has a purple arrow pointing to it.
- Loss function:**  $l(f(x^{(i)}; W), y^{(i)})$  is highlighted with a brown dashed box and has a brown arrow pointing to it.

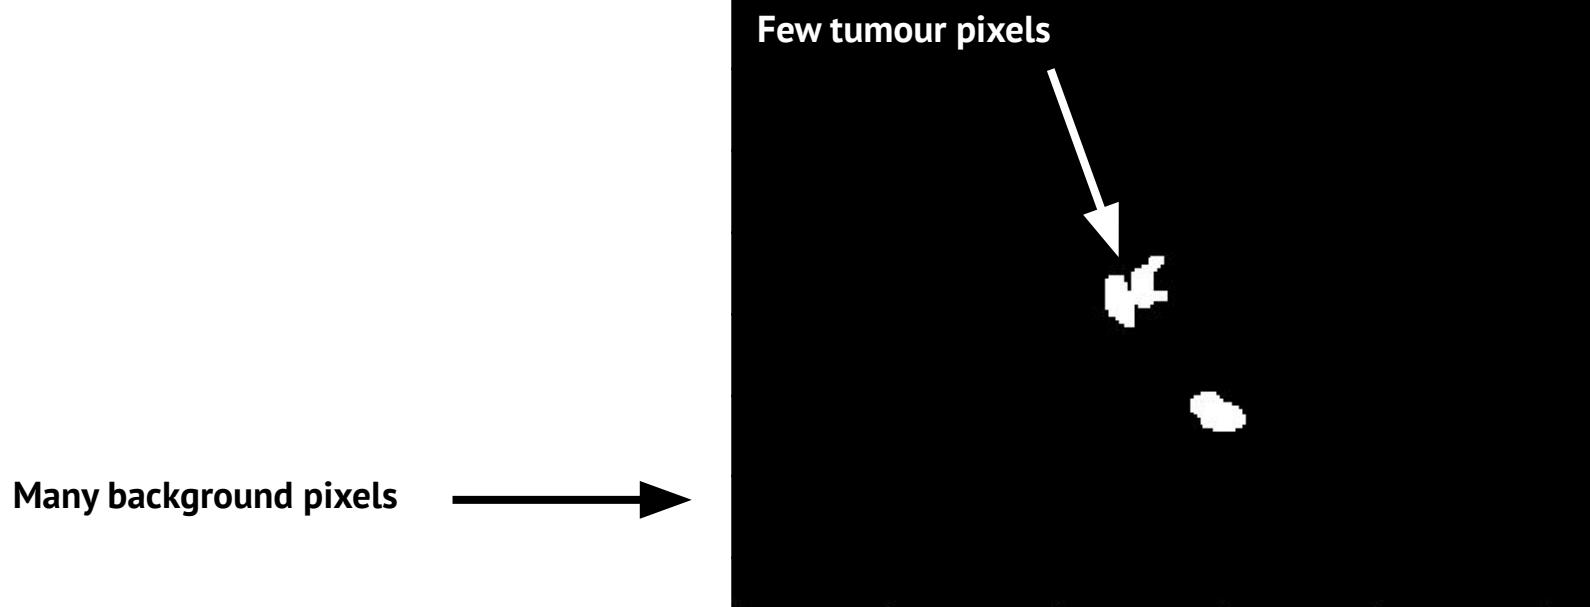
**For classification problems, we often use the cross entropy error used in logistic regression**

$$l_{\text{CE}}(f(x^{(i)}; W), y^{(i)}) = - \sum_p \left[ y_p^{(i)} \log \left( f \left( x_p^{(i)} \right) \right) + (1 - y_p^{(i)}) \log \left( 1 - f \left( x_p^{(i)} \right) \right) \right]$$

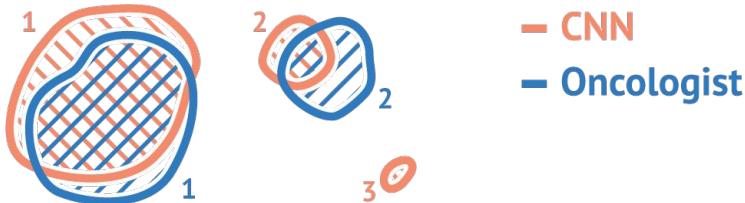
True pixel class

Predicted pixel class

**However, the cross-entropy loss puts too much weight into correctly classifying background pixels**



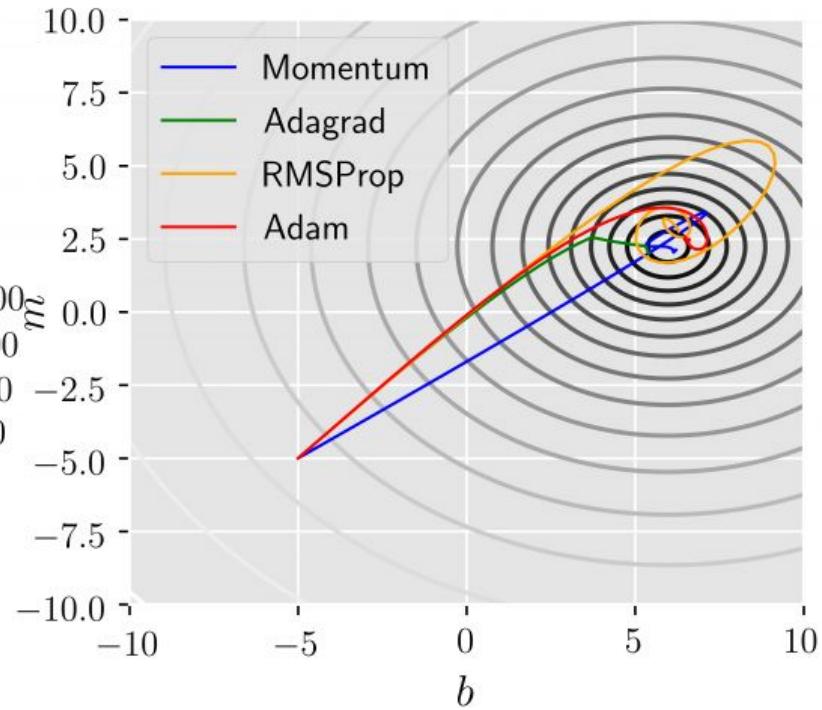
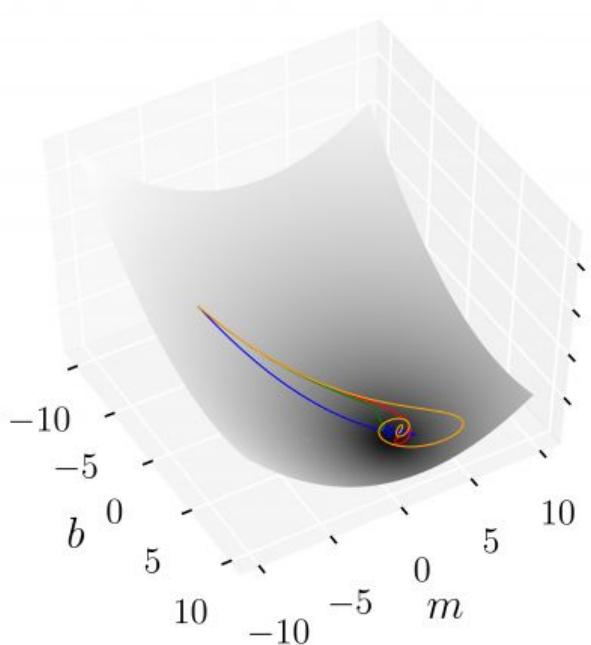
To combat this, we use a Dice loss function, that doesn't take the number of true negatives into account



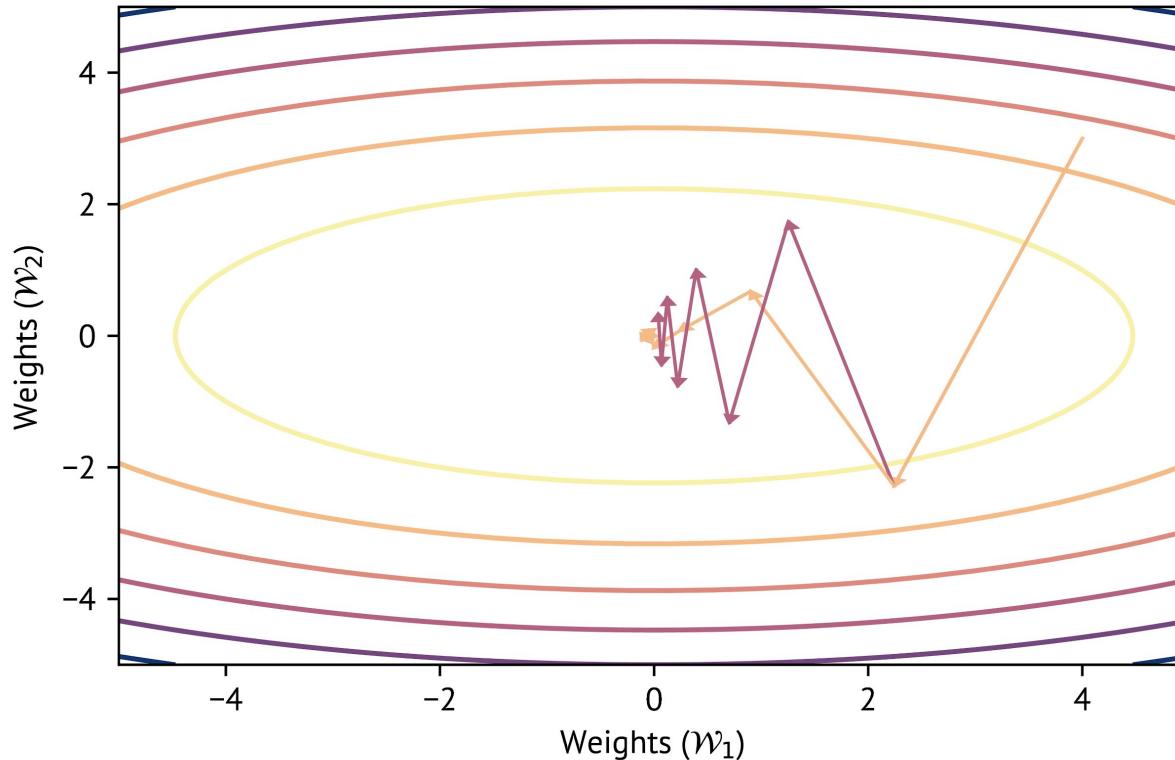
$$\text{Dice} = \frac{\frac{\square}{\square + \square}}{2} = 0.7$$

$$l_{\text{dice}}(f(x^{(i)}; W), y^{(i)}) = \frac{2 \sum_p \left[ f(x_p^{(i)}) y_p^{(i)} \right]}{\sum_p f(x_p^{(i)}) + \sum_p y_p^{(i)}}$$

**First minimise the loss using an adaptive momentum method such as Adam**



To minimise the loss, we use a stochastic first order optimisation method such as ADAM



If the generalisation error is large, try using momentum SGD optimiser instead of an adaptive momentum method

